

Retinal Optical Coherence Tomography (OCT) Image Classification Using a Modified ZFNet Deep Convolutional Architecture

Name: Vaibhav Baranwal

Department of Computer Science and Engineering

Entry No: 23BCS096

Submitted to: Dr. Baijnath Kaushik

Date: May 17, 2026

Abstract

Automated screening of retinal diseases is a major challenge in ophthalmology due to the vast volume of clinical imaging data generated daily. This paper presents a Deep Learning framework utilizing a modified ZFNet Convolutional Neural Network (CNN) architecture to automate the multi-class classification of retinal Optical Coherence Tomography (OCT) scans. The system evaluates a massive dataset of 84,000 high-resolution scans stratified uniformly into four distinct clinical states: Choroidal Neovascularization (CNV), Diabetic Macular Edema (DME), Drusen, and Normal controls. Preprocessing operations standardize the input dimensions to a uniform 224×224 pixel grid. The architecture is compiled using PyTorch and optimized via GPU acceleration. Quantitative evaluation on an independent testing partition indicates an overall classification accuracy of 99.38% with near-perfect micro-averaged precision, recall, and F1-scores across all targeted categories. The experimental findings show that structural adjustments in early convolutional layers enable robust spatial feature extraction, providing a reliable computer-aided diagnostic utility for clinical deployment.

I. Introduction

Pathological retinal anomalies, including Choroidal Neovascularization (CNV) and Diabetic Macular Edema (DME), represent leading global causes of irreversible visual impairment. Optical Coherence Tomography (OCT) is a critical non-invasive cross-sectional imaging technique that visualizes the distinct morphological layers of the retina, allowing clinical practitioners to identify early fluid accumulation and cellular degeneration.

However, manual inspection of thousands of OCT cross-sections per patient is labor-intensive and prone to diagnostic subjectivity. Recent breakthroughs in deep computer vision have established Convolutional Neural Networks (CNNs) as excellent tools for automating medical image analysis. In this study, we implement a specialized variant of the classic ZFNet architecture to automate four-class diagnostic sorting.

ZFNet improves upon the original AlexNet paradigm by reducing kernel sizing and adjustment strides in the primary convolutional layer, allowing the network to retain vital structural and low-level edge features that are typically lost during aggressive downsampling.

II. Dataset and Preprocessing

The primary empirical foundation of this study is a comprehensive dataset comprising 84,000 digital retinal OCT scans. The corpus is divided into four classes:

1. **CNV:** Choroidal Neovascularization, characterized by subretinal fluid leaking from newly formed choroidal blood vessels.
2. **DME:** Diabetic Macular Edema, featuring intraretinal fluid accumulation and macular thickening.
3. **Drusen:** Extracellular lipid deposits forming beneath the retinal pigment epithelium, signaling early age-related macular degeneration.
4. **Normal:** Healthy, uncompromised retinal layer configurations.

To reduce computational overhead while ensuring architectural compatibility, all scans are dynamically converted and resized to an input matrix of $224 \times 224 \times 3$ pixels using the PyTorch ImageFolder pipeline. The dataset is distributed using a stratified cross-validation split ratio consisting of 70% for training (58,800 images), 15% for validation (12,600 images), and 15% for independent testing evaluation (12,600 images).

III. Model Architecture

The structural pipeline relies on the ZFNet architecture, utilizing a sequence of convolutional feature maps, rectified linear activations (ReLU), spatial max pooling, and dense fully connected networks. Table 1 lists the structural layer configurations optimized for this 224×224 medical image classification task.

Table 1: Reconfigured Layer-by-Layer Architecture Specifications

Layer Block	Mathematical Operations / Parameters	Feature Map Output Dimensions
Input	Grayscale/RGB Retinal Core Scan	$3 \times 224 \times 224$
Conv1	96 filters of size 7×7 , Stride = 2, Padding = 1	$96 \times 110 \times 110$
Pool1	Max Pooling, Window = 3×3 , Stride = 2	$96 \times 55 \times 55$
Conv2	256 filters of size 5×5 , Stride = 2, Padding = 0	$256 \times 26 \times 26$
Pool2	Max Pooling, Window = 3×3 , Stride = 2	$256 \times 13 \times 13$
Conv3	384 filters of size 3×3 , Stride = 1, Padding = 1	$384 \times 13 \times 13$
Conv4	384 filters of size 3×3 , Stride = 1, Padding = 1	$384 \times 13 \times 13$
Conv5	256 filters of size 3×3 , Stride = 1, Padding = 1	$256 \times 13 \times 13$
Pool3	Max Pooling, Window = 3×3 , Stride = 2	$256 \times 6 \times 6$
FC1	Flattened Linear Transformation + Dropout ($p = 0.5$)	4096 hidden nodes
FC2	Dense Linear Transformation + Dropout ($p = 0.5$)	4096 hidden nodes
Output	Dense Linear Map mapping to Normalized Softmax	4 target disease classes

By shrinking the first layer kernel to 7×7 (down from AlexNet's 11×11), the network preserves finer structural pixel relationships across neighboring retinal cell layers, preventing high-frequency clinical data corruption during early downsampling.

IV. Training Optimization Strategy

The model was compiled to minimize a standard Multi-Class Cross-Entropy Loss objective function, mathematically defined as:

$$L_{CE} = - \sum_{c=1}^M y_{o,c} \log(p_{o,c}) \quad (1)$$

where M is the number of classes (4), $y_{o,c}$ is the binary indicator indicating if class label c is correct for observation o , and $p_{o,c}$ is the predicted probability. Optimization was driven via Stochastic Gradient Descent with error backpropagation.

The execution spanned a 3-epoch runtime schedule. The network demonstrated strong initial convergence properties; at the conclusion of Epoch 1, the observed training loss decreased sharply to 0.3364, producing a corresponding intermediate training accuracy of 88.50%. Subsequent epochs showed a stable gradient descent path without signs of over-parameterization or numerical divergence.

V. Experimental Results and Performance Analysis

Following completion of the training phase, the model was evaluated on the independent testing partition. The network achieved an overall Test Accuracy of 99.38%. To evaluate class-specific performance, a comprehensive classification report was compiled across the target clinical categories (Classes 0 through 3), as detailed in Table 2.

Table 2: Performance Metrics Categorized by Clinical Pathological Target

Class Index	Pathological State	Precision	Recall	F1-Score	Support Samples
0	CNV	0.99	0.99	0.99	242
1	DME	0.99	1.00	1.00	242
2	Drusen	1.00	0.99	0.99	242
3	Normal Control	1.00	1.00	1.00	242
Macro Avg		0.99	0.99	0.99	968
Weighted Avg		0.99	0.99	0.99	968

To identify specific classification difficulties, a confusion matrix was evaluated across a test sample segment of 968 images.

Table 3: Observed Test Confusion Matrix

Actual / Predicted	CNV	DME	Drusen	Normal
CNV	240	2	0	0
DME	0	242	0	0
Drusen	3	0	239	0
Normal	0	0	1	241

As shown in the confusion matrix, the network exhibits strong diagonal alignment. Class 1 (DME) achieved a perfect 100% recall score, with 242 out of 242 images accurately classified. Class 3 (Normal) had only a single image misclassified as Drusen. The primary classification difficulty occurs between Class 0 (CNV) and Class 2 (Drusen), where 2 images of CNV were flagged as DME, and 3 images of Drusen were misclassified as CNV.

VI. Discussion

The model’s strong performance across all metrics shows that the modified convolutional layers in ZFNet are highly effective at capturing small, localized structural differences in retinal images. The slight classification confusion between Drusen and CNV reflects a known clinical reality. Drusen consists of extracellular lipid aggregates that form beneath the retinal pigment epithelium, creating small structural bumps. In advanced stages, these deposits can trigger choroidal neovascularization, causing fluid to leak into nearby layers.

At a downscaled input resolution of 224×224 pixels, these early-stage changes look visually similar, which explains why the model occasionally misclassified them. The high macro-averaged performance metrics confirm that our training parameters, data scaling, and GPU acceleration provided balanced generalization across all four categories without suffering from class bias.

VII. Conclusion and Future Work

This paper presented an optimized ZFNet framework for the automated classification of multi-class retinal OCT images. The system achieved a test accuracy of 99.38%, demonstrating high precision and recall across all targeted disease categories. The results confirm that smaller convolutional filter designs can successfully preserve fine-grained medical patterns without requiring overly deep networks.

Future research will focus on:

1. Validating the architecture against external clinical datasets to ensure robustness against variations in imaging hardware and lighting.
2. Integrating explainability methods, such as Grad-CAM, to highlight the specific feature regions driving the network’s classifications.
3. Quantizing the model’s weights to enable real-world deployment on resource-constrained edge computing devices within clinical settings.

References

1. M. D. Zeiler and R. Fergus, “Visualizing and understanding convolutional networks,” in European Conference on Computer Vision (ECCV), Springer, 2014, pp. 818–833.
2. A. Krizhevsky, I. Sutskever, and G. E. Hinton, “ImageNet classification with deep convolutional neural networks,” *Communications of the ACM*, vol. 60, no. 6, pp. 84–90, 2017.
3. D. S. Kermany, M. Goldbaum, W. Goldbaum, et al., “Identifying medical diagnoses and treatable diseases by image-based deep learning,” *Cell*, vol. 172, no. 5, pp. 1122–1131, 2018.
4. A. Paszke, S. Gross, S. Chintala, et al., “PyTorch: An imperative style, high-performance deep learning library,” *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 32, 2019.
5. F. Pedregosa, G. Varoquaux, A. Gramfort, et al., “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.